
Project Proposal: Efficient Exploration of LLM Agent Topologies with Graph Neural Networks

Gal Zmazek

University of Ljubljana
Ljubljana, Slovenia
gz66721@student.uni-lj.si

Vito Levstik

University of Ljubljana
Ljubljana, Slovenia
vl40930@student.uni-lj.si

1 Project Overview

The rapid advancement of large language models (LLMs) has enabled their widespread use across diverse domains. One promising direction is the multi-agent approach [5, 17], where multiple specialized LLM agents are interconnected in a graph-like structure to collaboratively solve complex problems. However, a major challenge lies in determining which graph topology of agent interactions leads to optimal performance [20]. Exhaustively testing all possible configurations is computationally and financially infeasible due to the cost of LLM evaluations [16].

To address this, the project aims to leverage graph neural networks (GNNs) to approximate and predict the collective performance of multi-agent LLM systems [19]. This approach seeks to efficiently explore the space of possible graph structures, providing insights into how network design influences reasoning quality and enabling more scalable and effective multi-agent LLM systems.

2 Application Domain

The project targets mathematical problem solving as the primary application domain. To support this, we will use publicly available datasets containing mathematical problems and their corresponding solutions [2, 7, 10]. The datasets consists of problem–solution pairs, which will be used to evaluate the performance of multi-agent large language structure. Evaluation will primarily focus on solution accuracy as the key performance metric for each multi-agent structure.

3 Main Pipeline

The proposed framework operates through an iterative pipeline with four interconnected components: the *Graph Evaluation Dataset*, *Graph Generation Block*, *GNN Block*, and *LLM Graph Evaluation Block*. Figure 1 illustrates the overall architecture. The pipeline progressively refines our understanding of which multi-agent graph topologies yield superior mathematical reasoning performance while minimizing expensive LLM evaluations through learned models.

Graph Evaluation Dataset The *Graph Evaluation Dataset* serves as the central knowledge repository, storing pairs of graph structures and their corresponding performance evaluations. Each entry consists of a graph G , paired with its empirically measured performance $p \in \mathbb{R}$. This dataset provides training data $\{(G_i, p_i)\}_{i=1}^N$ for the GNN model and grows iteratively as new evaluations are completed.

Graph Generation Block The *Graph Generation Block* produces candidate multi-agent topologies for evaluation. It takes the current *Graph Evaluation Dataset* as input and outputs a batch of K candidate graph structures. Initially, graphs are sampled randomly from a constrained space. As the dataset grows, the generator transitions to informed sampling strategies including similarity-based generation from top-performing graphs, using graph similarity algorithms [3, 8, 12].




Figure 1: Schematic representation of pipeline.

GNN Block The *GNN Block* functions as a surrogate model that predicts multi-agent system performance without requiring costly LLM evaluations. It takes as input the *Graph Evaluation Dataset* $\{(G_i, p_i)\}_{i=1}^N$ for training and K unevaluated candidate graphs from the generation block. It outputs a trained GNN model $f_\theta : \mathcal{G} \rightarrow \mathbb{R}$, predicted performance values $\{\hat{p}_1, \dots, \hat{p}_K\}$, and a filtered subset of top- M graphs (typically $M \ll K$). The task is therefore graph prediction.

The choice of GNN architecture is motivated by the strong dependency between agent interactions and their local neighborhood properties. In multi-agent reasoning systems, an agent’s contribution depends critically on the types and features of agents it communicates with. We will primarily focus on two attention-based architectures: Graph Attention Networks (GAT) [11] and Heterogeneous Graph Attention Networks (HetGAT) [13].

For comparison, we also evaluate simpler baseline models including Graph Convolutional Networks (GCN) [6], which perform uniform neighborhood averaging, and GraphSAGE [4], which uses sampling-based aggregation. These baselines allow us to quantify the benefit of attention mechanisms for this problem domain.

LLM Graph Evaluation Block The *LLM Graph Evaluation Block* performs empirical evaluation of selected multi-agent structures on mathematical reasoning tasks. It takes as input the top- M candidate graphs from the GNN Block and a fixed problem set $\mathcal{P} = \{(q_1, a_1), \dots, (q_T, a_T)\}$ from our Math problem dataset. It outputs performance scores $\{p_j\}$ and new graph-performance pairs $\{(G_j, p_j)\}_{j=1}^M$ that are appended to the Graph Evaluation Dataset.

The evaluation infrastructure is implemented in Python using LangGraph [1] for multi-agent orchestration and open-source LLM models. To establish performance baselines, we compare our learned multi-agent topologies against classical single-agent prompting strategies including Chain-of-Thought reasoning [15], Self-Consistency [14], Reflection [9], and Tree-of-Thoughts [18].

4 Hypotheses

We formulate three core hypotheses to guide our experimental investigation:

H1 As the Graph Evaluation Dataset expands through iterative evaluations, the GNN surrogate model’s predictive accuracy will monotonically improve. We will track correlation metric between predicted and actual performance scores to quantify this improvement over successive iterations.

H2 Different base LLM models (e.g., GPT-4, Claude, Llama) and LLM type of agent configurations will yield distinct optimal multi-agent graph structures.

H3 As the GNN improves, it will preferentially select high-performing candidates for evaluation, leading to over-representation of successful topologies in the dataset. This bias will cause the model to overestimate performance of new candidates, resulting in disappointing empirical evaluations. This negative feedback naturally re-calibrates the GNN, creating a self-correcting mechanism.

References

- [1] LangChain AI. Langgraph: Multi-agent workflows. <https://langchain-ai.github.io/langgraph/>, 2024. Accessed: 2025-10-21.
- [2] Aida Amini, Saadia Gabriel, Shanchuan Lin, Rik Koncel-Kedziorski, Yejin Choi, and Hannaneh Hajishirzi. MathQA: Towards interpretable math word problem solving with operation-based formalisms. In Jill Burstein, Christy Doran, and Thamar Solorio, editors, *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 2357–2367, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics.
- [3] Shri Prakash Dwivedi. Some algorithms on exact, approximate and error-tolerant graph matching, 2020.
- [4] William L. Hamilton, Rex Ying, and Jure Leskovec. Inductive representation learning on large graphs, 2018.
- [5] Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding, 2021.
- [6] Thomas N. Kipf and Max Welling. Semi-supervised classification with graph convolutional networks, 2017.
- [7] David Saxton, Edward Grefenstette, Felix Hill, and Pushmeet Kohli. Analysing mathematical reasoning abilities of neural models, 2019.
- [8] Nino Shervashidze, Pascal Schweitzer, Erik Jan van Leeuwen, Kurt Mehlhorn, and Karsten M. Borgwardt. Weisfeiler-lehman graph kernels. *Journal of Machine Learning Research*, 12(77):2539–2561, 2011.
- [9] Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. *Advances in Neural Information Processing Systems*, 36, 2023.
- [10] Shubham Toshniwal, Ivan Moshkov, Sean Narenthiran, Daria Gitman, Fei Jia, and Igor Gitman. Openmathinstruct-1: A 1.8 million math instruction tuning dataset, 2024.
- [11] Petar Veličković, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Liò, and Yoshua Bengio. Graph attention networks, 2018.
- [12] Mengzhao Wang, Xiaoliang Xu, Qiang Yue, and Yuxiang Wang. A comprehensive survey and experimental comparison of graph-based approximate nearest neighbor search, 2021.
- [13] Xiao Wang, Houye Ji, Chuan Shi, Bai Wang, Peng Cui, P. Yu, and Yanfang Ye. Heterogeneous graph attention network, 2021.
- [14] Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models. *arXiv preprint arXiv:2203.11171*, 2022.
- [15] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models, 2023.
- [16] Jiayi Yang, Mengqi Zhang, Yiqiao Jin, Hao Chen, Qingsong Wen, Lu Lin, Yi He, Srikanth Kumar, Weijie Xu, James Evans, and Jindong Wang. Topological structure learning should be a research priority for llm-based multi-agent systems, 2025.
- [17] Sen Yang, Yafu Li, Wai Lam, and Yu Cheng. Multi-llm collaborative search for complex problem solving, 2025.
- [18] Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran, Thomas L Griffiths, Yuan Cao, and Karthik Narasimhan. Tree of thoughts: Deliberate problem solving with large language models. *Advances in Neural Information Processing Systems*, 36, 2023.

- [19] Guibin Zhang, Yanwei Yue, Xiangguo Sun, Guancheng Wan, Miao Yu, Junfeng Fang, Kun Wang, Tianlong Chen, and Dawei Cheng. G-designer: Architecting multi-agent communication topologies via graph neural networks, 2025.
- [20] Han Zhou, Xingchen Wan, Ruoxi Sun, Hamid Palangi, Shariq Iqbal, Ivan Vulić, Anna Korhonen, and Serkan Ö. Arık. Multi-agent design: Optimizing agents with better prompts and topologies, 2025.